

Coursework2__CST4050__

April 22, 2020

1 1. Introduction:

Yelp is a public company which publishes crowd-sourced reviews about local businesses, as well as the online reservation service Yelp Reservations. It also assists small businesses in how to respond to reviews. The company was founded in the year 2004 by former PayPal employees. By the year 2010 it had \$30 million in revenues and the website had published more than 4.5 million crowd-sourced reviews.

2 2. Dataset Description

The data used here is collected from Yelp website (<https://www.yelp.com/dataset>). The data was uploaded there for learning/academic purposes. The data consists of many JSON files. There are 6 JSON files available on the website. For my project I have used 2 of them. One file consists of the business data. It contains many vital business data. The other file used in the analysis consists of the reviews received from the users. The details of the attributes of these two files are listed below:

Business Data:

"business_id": It is a unique string business id.

"name": It is a string character. It is name of the business.

"address": It is a string character, It is the full address of the business.

"city": It is a string character, It consists of the city where the business is located.

"state": It is a string character, It consists of the character state code, if applicable.

"postal code": It is a string character, It consists of the postal code of the location of the business.

"latitude": It is a float character. It is the latitude of the location.

"longitude": It is a float character. It is the longitude of the location.

"stars": It is a float character. It consists of the star rating, rounded to half-stars.

"review_count": It is an integer character, it consists of a number of reviews.

"is_open": It is an integer, 0 or 1 for closed or open, respectively.

"attributes": It is an object, business attributes to values. note: some attribute values might be objects.

"categories": It is an array of strings of business categories.

"hours": It is an object of key day to value hours, hours are using a 24hr clock.

Review Data:

"review_id": It uses a string character. It is an unique review id.

"user_id": It is a string character. It is an unique user id that maps to the user in user.json.

"business_id": It is a string character. It is a business id that maps to business in business.json.

"stars": It is an integer that consists of the star rating.

"date": It is a string character which is in date format YYYY-MM-DD.

"text": It is a string character which consists of the reviews itself.

"useful": It is an integer which consists of the number of useful votes received.

"funny": It is an integer which consists of the number of funny votes received.

"cool": It is an integer which consists of number of cool votes received.

```
[231]: # Loading the required libraries

# numpy is the library imported for doing the linear algebra
import numpy as np

# pandas are used for data processing, JSON file I/O (e.g. pd.read_JSON)
import pandas as pd

import collections # this is used to store collections of data
import re, string # this is used for string searching and manipulation
import sys # the sys module provides information about constants, functions and
    ↪ methods of python interpreter.
import time
```

```
[232]: # The below function loads the JSON file and converts it into a dataframe in
    ↪ pandas.
import json

def init_ds(json):
    ds= {}
    keys = json.keys()
    for k in keys:
        ds[k]= []
    return ds, keys

def read_json(file):
    dataset = {}
    keys = []
    with open(file,encoding="utf8") as file_lines:
```

```

for count, line in enumerate(file_lines):
    data = json.loads(line.strip())
    if count == 0:
        dataset, keys = init_ds(data)
    for k in keys:
        dataset[k].append(data[k])

return pd.DataFrame(dataset)

```

```

[3]: yelp_business= read_json('yelp_academic_dataset_business.json') # The business_
    ↪ data file
yelp_review= read_json('yelp_academic_dataset_review.json') # The review data_
    ↪ file

```

First glance at the review file

```
[5]: yelp_review.head()
```

```

[5]:
      review_id      user_id      business_id \
0  xQY8N_XvtGbearJ5X4QryQ  OwjRMXRCOKyPrIlcjaXeFQ  -MhfebMOQIsKt87iDN-FNw
1  UmFMZ8PyXZTY2QcwzsfQYA  nIJD_7ZXHq-FX8byPMOkMQ  1brU8StCq3yDfr-QMnGrmQ
2  LG2ZaYiOgpr2DK_90pYjNw  V34qejsxNsCbcgD8COHVk-Q  HQ128KMwrEKHqhFrrDqVNQ
3  i6g_oA9Yf9Y31qt0wibXpw  ofKDKJKXSKZXu5xJNGiiBQ  5JxlZaqCnk1MnbgRirs40Q
4  6TdNDKywdbjoTkizeMce8A  UgMW8bLE0QMJDCKQ1Ax5Mg  IS4cv902ykd8wj1TRON3-A

      stars  useful  funny  cool  \
0      2.0      5      0      0
1      1.0      1      1      0
2      5.0      1      0      0
3      1.0      0      0      0
4      4.0      0      0      0

                                text      date
0  As someone who has worked with many museums, I...  2015-04-15 05:21:16
1  I am actually horrified this place is still in...  2013-12-07 03:16:52
2  I love Deagan's. I do. I really do. The atmo...  2015-12-05 03:18:11
3  Dismal, lukewarm, defrosted-tasting "TexMex" g...  2011-05-27 05:30:52
4  Oh happy day, finally have a Canes near my cas...  2017-01-14 21:56:57

```

First glance at the business file

```
[6]: yelp_business.head()
```

```

[6]:
      business_id      name \
0  f9NumwFMBDn751xgFiRbNA  The Range At Lake Norman
1  YzvJg0SayhoZgCljUJRF9Q      Carlos Santo, NMD
2  XNoUzKckATkOD1hP6vghZg      Felinus
3  60AZjbxqM5ol29BuHsil3w      Nevada House of Hose

```

```
4 51M2Kk903DFYI6gnB5I6SQ USE MY GUY SERVICES LLC
```

```

                                address          city state postal_code  latitude \
0          10913 Bailey Rd          Cornelius    NC          28031 35.462724
1 8880 E Via Linda, Ste 107          Scottsdale    AZ          85258 33.569404
2      3554 Rue Notre-Dame O          Montreal    QC          H4C 1P4 45.479984
3      1015 Sharp Cir North Las Vegas    NV          89030 36.219728
4      4827 E Downing Cir          Mesa        AZ          85205 33.428065

```

```

    longitude stars review_count is_open \
0 -80.852612   3.5           36        1
1 -111.890264   5.0           4         1
2 -73.580070   5.0           5         1
3 -115.127725   2.5           3         0
4 -111.726648   4.5          26         1

```

```

                                attributes \
0 {'BusinessAcceptsCreditCards': 'True', 'BikePa...
1 {'GoodForKids': 'True', 'ByAppointmentOnly': '...
2                                     None
3 {'BusinessAcceptsCreditCards': 'True', 'ByAppo...
4 {'BusinessAcceptsCreditCards': 'True', 'ByAppo...

```

```

                                categories \
0 Active Life, Gun/Rifle Ranges, Guns & Ammo, Sh...
1 Health & Medical, Fitness & Instruction, Yoga,...
2                                     Pets, Pet Services, Pet Groomers
3 Hardware Stores, Home Services, Building Suppl...
4 Home Services, Plumbing, Electricians, Handyma...

```

```

                                hours
0 {'Monday': '10:0-18:0', 'Tuesday': '11:0-20:0'...
1                                     None
2                                     None
3 {'Monday': '7:0-16:0', 'Tuesday': '7:0-16:0', ...
4 {'Monday': '0:0-0:0', 'Tuesday': '9:0-16:0', '...

```

3 3. Machine learning challenge

The challenge here is to build a model to classify the Yelp Reviews into 1 star, 3 star or 5 star(Negative,average,good) categories based on the text content.

We can also use the model for Sentiment Analysis and Prediction of Review Ratings.

This can be achieved by performing machine learning for "textual data analysis" . This allows us to extract and classify the reviews to make better predictions and create insights.

4 4. Better Understanding of the data

Before preparing the required dataset, I want to analyse what will be the best parameter to filter the data.

5 -- Top reviewed business(by name)

```
[233]: top_reviewed = yelp_review[yelp_review["stars"]>3]
top_reviews_dict = {}

for business_id in top_reviewed["business_id"].values:
    try :
        top_reviews_dict[business_id] =top_reviews_dict[business_id]+1
    except:
        top_reviews_dict[business_id]=1


topbusiness = pd.DataFrame.from_dict(data= top_reviews_dict,orient="index")

topbusiness.reset_index(inplace=True)
topbusiness.columns = ['business_id', 'rated']
del(top_reviews_dict)
del(top_reviewed)

[234]: top_count= 20
right=pd.DataFrame(yelp_business[['business_id',"name","categories"]].values,
                  columns=['business_id',"Business name","categories"])

top_business_data = pd.merge(topbusiness,right=right,
    ↪how="inner",on='business_id')
top_business_data.sort_values("rated")[::-1][:top_count].plot(x="Business_
    ↪name",y="rated",
                                                                    kind="bar",figsize=(14,6),
                                                                    title='Positive reviews').
    ↪set_ylabel("Total ratings")

del(topbusiness)
del(right)
```



The top reviewed business merchants are Mon Ami Gabi, Bacchanal Buffet, Hash House A Go Go. All these are restaurants.


6 -- Categories of top reviewed businesses

```
[9]: num_cat = 10 # to show top 10 catrgories
top_business = 30 # choose categories of top 30 businesses
cat_data = top_business_data.sort_values("rated")[::-1][:top_business]
#cat_data.categories
Categories={}
for cat in cat_data.categories.values:
    all_categories= cat.split(",")
    for x in all_categories:
        try :
            Categories[x] =Categories[x]+1
        except:
            Categories[x]=1
top_categories = pd.DataFrame.from_dict(data= Categories,orient="index")
top_categories.reset_index(inplace=True)
top_categories.columns = ['category', 'occurance']

x_val=top_categories.sort_values("occurance")[::-1][:num_cat].occurance.values
labels=top_categories.sort_values("occurance")[::-1][:num_cat].category.values
```

```
series = pd.Series(x_val, index=labels, name='Top business types')
series.plot.pie(figsize=(10, 10),startangle=90)
```

[9]: <matplotlib.axes._subplots.AxesSubplot at 0x16ad62ff0c8>



Restuarants are the top reviwed business categories.

7 -- Negatively reviewed businesses

```
[10]: bottom_reviewed = yelp_review[yelp_review["stars"]<2]
bottom_reviews_dict ={}

for business_id in bottom_reviewed["business_id"].values:
    try :
        bottom_reviews_dict[business_id] =bottom_reviews_dict[business_id]+1
    except:
```

```

        bottom_reviews_dict[business_id]=1

bottombusiness = pd.DataFrame.from_dict(data=bottom_reviews_dict,orient="index")

bottombusiness.reset_index(inplace=True)
#bottombusiness.head()
bottombusiness.columns = ['business_id', 'rated']

```


```

[11]: top_count= 20
right=pd.DataFrame(yelp_business[['business_id',"name","categories"]].values,
                    columns=['business_id',"Business name","categories"])

bottom_business_data = pd.merge(bottombusiness,right=right,
                                how="inner",on='business_id')
bottom_business_data.sort_values("rated")[::-1][:top_count].plot(x="Business_
                                name",y="rated",
                                kind="bar",figsize=(14,6),
                                title='Negative reviews').
                                set_ylabel("Total 1 star ratings")

del(bottom_reviewed)
del(bottom_reviews_dict)
del(bottombusiness)
del(right)

```

Casinos are the negatively reviewed business.

8 5. Methodology

9 5.1 Data Collection

From the above analysis, it is clear that the restaurants are the top reviewed business. Therefore, for sentimental analysis I want to collect the reviews for "Indian" restaurants only.

I will merge the two datafiles now. I want to develop the rating predictive model for the "Indian" restaurants only based on their reviews.

Since I want to collect the reviews only for the "Indian restaurants", I have to extract the category column as individual elements. So that I can filter my data accordingly. Further I want to collect the data for the Indian restaurants which are still open.

```
[12]: # The explode() function is used to transform each element of a list-like to a row.
```

```
# To use this function, pandas version needs to be above 0.25.
# Checking the version of pandas.
pd.__version__
```

```
[12]: '0.25.1'
```

The explode() function is used to transform each element of a list-like to a row

```
[13]: # Applying the explode function to the column "categories"
df_explode = yelp_business.assign(categories = yelp_business.categories
                                   .str.split(', ')).explode('categories')
```

```
[14]: #We can then list out all the individual category
df_explode.categories.value_counts()
```

```
[14]: Restaurants      63944
      Shopping         34644
      Food             32991
      Home Services    22487
      Beauty & Spas     20520
      ...
      Halfway Houses      1
      Street Art          1
      Osteopaths          1
      Tonkatsu            1
      Bungee Jumping      1
      Name: categories, Length: 1336, dtype: int64
```

I want to get the reviews for the Indian Restaurants

```
[15]: #Find the categories containing Indian Restaurants
df_explode[df_explode.categories.str.contains('Indian',
                                              case=True,na=False)].categories.value_counts()
```

```
[15]: Indian      1652
      Name: categories, dtype: int64
```

```
[16]: print('Total number of categories',len(df_explode.categories.value_counts()))
```

Total number of categories 1336

```
[17]: # Finding the categories that contains Restaurants
df_explode[df_explode['categories'].str.contains('Indian',case=True,na=False)].
    ↪categories.value_counts()
```

```
[17]: Indian      1652
      Name: categories, dtype: int64
```

```
[18]: # Extracting the data only for the Indian Restaurants from the business file.
business_Restaurants = yelp_business[yelp_business['categories'].str.contains(
    'Indian',case=False, na=False)]
```

```
[19]: business_Restaurants.head()
```

```
[19]:
```

	business_id	name \
82	RrapAhd8ZxCj-iue7fu9FA	Ganga Restaurant
218	Yy_SB9r7VQTlIVv-MeOmaA	Cari Mela
366	LZWXP-D4YPlzsFjVx6b9XA	Spice South
511	CvOxwAhmwwy3sQRKYVAjGg	Bombay Buffet Indian Cuisine
525	IjYGSm4_DjP6VULpP27o_g	My Roti Place

	address	city	state	postal_code	latitude \
82	515 4th Avenue SW	Calgary	AB	T2P 0J8	51.049407
218	2556 Rue Centre	Montréal	QC	H3K 1J8	45.478639
366	8145 Ardrey Kell Rd	Charlotte	NC	28277	35.039010
511	795 Markham Road	Scarborough	ON	M1H 2Y1	43.767071
525	948 Queen Street E	Toronto	ON	M4M 1J7	43.660961

	longitude	stars	review_count	is_open \
82	-114.072656	1.5	3	1
218	-73.568839	3.5	9	1
366	-80.793656	3.5	20	0
511	-79.228062	2.5	22	0
525	-79.341156	4.0	8	1

	attributes \
82	{'RestaurantsReservations': 'True', 'WiFi': 'u...
218	{'BusinessParking': {'garage': False, 'street...
366	{'OutdoorSeating': 'True', 'RestaurantsGoodFor...
511	{'Caters': 'False', 'RestaurantsPriceRange2': ...
525	{'OutdoorSeating': 'True', 'GoodForMeal': {'d...

	categories \
82	Restaurants, Buffets, Indian, Halal
218	Indian, Restaurants
366	Indian, Restaurants
511	Restaurants, Indian, Event Planning & Services...
525	Indian, Restaurants, Fast Food

	hours
82	{'Monday': '17:0-22:0', 'Tuesday': '17:0-22:0'...
218	{'Monday': '11:30-22:0', 'Tuesday': '11:30-22:...
366	{'Tuesday': '11:30-21:30', 'Wednesday': '11:30...
511	None
525	{'Monday': '11:0-21:0', 'Tuesday': '11:0-21:0'...

```
[20]: # Collecting the data only for the Indian restaurants which are still open.
# 1 = open, 0 = closed
business_Restaurants = business_Restaurants[business_Restaurants['is_open']==1]
```

```
[21]: # Remove the columns that are not required in business_Restaurants
drop_columns = ['city', 'state', 'postal_code', 'latitude', 'longitude', 'hours']
business_Restaurants = business_Restaurants.drop(drop_columns, axis=1)
business_Restaurants.head()
```

```
[21]:
```

	business_id	name \
82	RrapAhd8ZxCj-iue7fu9FA	Ganga Restaurant
218	Yy_SB9r7VQTlIVv_MeOmaA	Cari Mela
525	IjYGSm4_DjP6VULpP27o_g	My Roti Place
650	ZDx7kt4aOPTlmYTqXDrTGA	Canbe Foods
969	RE4qn28MEiDrM1PbdYVgxA	Lena's Roti & Doubles

	address	stars	review_count	is_open \
82	515 4th Avenue SW	1.5	3	1
218	2556 Rue Centre	3.5	9	1
525	948 Queen Street E	4.0	8	1
650	336 Rossland Road E	4.0	15	1
969	100 Maritime Ontario Boulevard, Unit 69	4.0	20	1


```

attributes \
82 {'RestaurantsReservations': 'True', 'WiFi': 'u...
218 {'BusinessParking': {'garage': False, 'street...
525 {'OutdoorSeating': 'True', 'GoodForMeal': {'d...
650 {'Caters': 'True', 'WheelchairAccessible': 'Tr...
969 {'RestaurantsGoodForGroups': 'True', 'GoodForK...

categories
82 Restaurants, Buffets, Indian, Halal
218 Indian, Restaurants
525 Indian, Restaurants, Fast Food
650 Indian, Restaurants, Sri Lankan
969 Caribbean, Restaurants, Indian
```

Review data file is a huge file. If we try to load all the data at once, it is likely to crash the memory of the computer. Therefore we will load large data by segmenting the file into smaller chunks. Also to reduce the memory usage I am identifying the datatype of each column.

```
[22]: # To reduce the memory usage identifying the datatype of each column
size = 100000
review = pd.read_json('yelp_academic_dataset_review.json', lines=True,
dtype={'review_id':str, 'user_id':str,
'business_id':str, 'stars':int,
'date':str, 'text':str, 'useful':int,
```

```
        'funny':int,'cool':int},
        chunksize=size)
```

```
[23]: # There are multiple chunks to be read
chunk_list = []
for chunk_review in review:
    # Drop columns that aren't needed
    chunk_review = chunk_review.drop(['review_id','date'], axis=1)
    # Renaming column name to avoid conflict with business overall star rating
    chunk_review = chunk_review.rename(columns={'stars': 'review_stars'})
    # Inner merge with edited business file so only reviews related to the
    ↪business remain
    chunk_merged = pd.merge(business_Restaurants, chunk_review,
    ↪on='business_id', how='inner')
    # Show feedback on progress
    print(f"{chunk_merged.shape[0]} out of {size:,} related reviews")
    chunk_list.append(chunk_merged)
# After trimming down the review file, concatenate all relevant data back to one
    ↪dataframe
df = pd.concat(chunk_list, ignore_index=True, join='outer', axis=0)
```

```
787 out of 100,000 related reviews
780 out of 100,000 related reviews
800 out of 100,000 related reviews
793 out of 100,000 related reviews
810 out of 100,000 related reviews
862 out of 100,000 related reviews
867 out of 100,000 related reviews
905 out of 100,000 related reviews
892 out of 100,000 related reviews
957 out of 100,000 related reviews
934 out of 100,000 related reviews
838 out of 100,000 related reviews
814 out of 100,000 related reviews
778 out of 100,000 related reviews
801 out of 100,000 related reviews
958 out of 100,000 related reviews
1464 out of 100,000 related reviews
1430 out of 100,000 related reviews
1419 out of 100,000 related reviews
1460 out of 100,000 related reviews
1298 out of 100,000 related reviews
1105 out of 100,000 related reviews
985 out of 100,000 related reviews
964 out of 100,000 related reviews
1014 out of 100,000 related reviews
1082 out of 100,000 related reviews
```

1049 out of 100,000 related reviews
1047 out of 100,000 related reviews
1009 out of 100,000 related reviews
1083 out of 100,000 related reviews
1071 out of 100,000 related reviews
1061 out of 100,000 related reviews
670 out of 100,000 related reviews
580 out of 100,000 related reviews
642 out of 100,000 related reviews
591 out of 100,000 related reviews
586 out of 100,000 related reviews
681 out of 100,000 related reviews
654 out of 100,000 related reviews
593 out of 100,000 related reviews
910 out of 100,000 related reviews
1109 out of 100,000 related reviews
1083 out of 100,000 related reviews
1036 out of 100,000 related reviews
1035 out of 100,000 related reviews
969 out of 100,000 related reviews
856 out of 100,000 related reviews
783 out of 100,000 related reviews
738 out of 100,000 related reviews
692 out of 100,000 related reviews
734 out of 100,000 related reviews
722 out of 100,000 related reviews
723 out of 100,000 related reviews
773 out of 100,000 related reviews
881 out of 100,000 related reviews
860 out of 100,000 related reviews
1081 out of 100,000 related reviews
1055 out of 100,000 related reviews
1082 out of 100,000 related reviews
1134 out of 100,000 related reviews
1039 out of 100,000 related reviews
1033 out of 100,000 related reviews
935 out of 100,000 related reviews
895 out of 100,000 related reviews
922 out of 100,000 related reviews
901 out of 100,000 related reviews
897 out of 100,000 related reviews
905 out of 100,000 related reviews
933 out of 100,000 related reviews
943 out of 100,000 related reviews
834 out of 100,000 related reviews
840 out of 100,000 related reviews
1168 out of 100,000 related reviews
1284 out of 100,000 related reviews

1236 out of 100,000 related reviews
 1246 out of 100,000 related reviews
 1139 out of 100,000 related reviews
 1133 out of 100,000 related reviews
 823 out of 100,000 related reviews
 863 out of 100,000 related reviews
 178 out of 100,000 related reviews

We have finally collected our required data for the application of our machine learning models (Sentimental analysis)..

```
[24]: # To view the collected data
df.head()
```

```
[24]:
```

	business_id	name \
0	RrapAhd8ZxCj-iue7fu9FA	Ganga Restaurant
1	Yy_SB9r7VQTlIVv_MeOmaA	Cari Mela
2	ZDx7kt4aOPTlmYTqXDrTGA	Canbe Foods
3	RE4qn28MEiDrM1PbdYVgxA	Lena's Roti & Doubles
4	_PKXarw3GjlbwbXhjdpUMA	Tangra Villa Hakka Chinese Cuisine

	address	stars	review_count	is_open \
0	515 4th Avenue SW	1.5	3	1
1	2556 Rue Centre	3.5	9	1
2	336 Rossland Road E	4.0	15	1
3	100 Maritime Ontario Boulevard, Unit 69	4.0	20	1
4	411 Manhattan Drive, Suite 3C	3.5	29	1


```

attributes \
0 {'RestaurantsReservations': 'True', 'WiFi': 'u...
1 {'BusinessParking': {'garage': False, 'street...
2 {'Caters': 'True', 'WheelchairAccessible': 'Tr...
3 {'RestaurantsGoodForGroups': 'True', 'GoodForK...
4 {'RestaurantsAttire': 'u'casual'', 'Caters': '...'

```


	categories	user_id	review_stars \
0	Restaurants, Buffets, Indian, Halal	pBUsRjJLTN-TVsoIv5ue8w	3
1	Indian, Restaurants	h1UcaSPIPpQqMiWf12Csqw	4
2	Indian, Restaurants, Sri Lankan	zFGpxwJewI60jC2u9EnZ-g	5
3	Caribbean, Restaurants, Indian	F-nLfSJ4qtdZRwNW993lSw	4
4	Chinese, Indian, Restaurants	R_TJQ6Hy1BtOYK_hC125bQ	5

	useful	funny	cool	text
0	0	0	0	Ordered biryani for tonight's supper. I wasn't...
1	2	1	1	We started with a mix platter which had all ki...
2	7	1	2	You know this place is good when most of the p...
3	0	0	0	Visited this place with my Office colleagues, ...
4	1	1	0	So spicy and flavorful! Shrimp Pakoras really ...

10 5.2 Data Analysis Process

11 5.2.1 Data Preprocessing

The preprocessing of the data includes removing the unwanted columns from the dataset. It also includes eliminating stop words and punctuations. This is achieved by using nltk in python. It stands for natural language Toolkit. It is a suite of libraries and programs for symbolic and statistical natural language processing. I have also created a new column in the name as "text length". It returns the number of words in the column "text". Since, text data requires special preparation before we start it for predictive modeling, text must be parsed to remove words, called tokenization. For the input to a machine learning algorithm, the words need to be encoded as integers or floating point values. This is called feature extraction (or vectorization).

```
[25]: # To get the information about the object types in the new data.  
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 75517 entries, 0 to 75516  
Data columns (total 14 columns):  
business_id      75517 non-null object  
name             75517 non-null object  
address         75517 non-null object  
stars            75517 non-null float64  
review_count     75517 non-null int64  
is_open         75517 non-null int64  
attributes       75375 non-null object  
categories       75517 non-null object  
user_id         75517 non-null object  
review_stars     75517 non-null int32  
useful          75517 non-null int32  
funny           75517 non-null int32  
cool            75517 non-null int32  
text            75517 non-null object  
dtypes: float64(1), int32(4), int64(2), object(7)  
memory usage: 6.9+ MB
```

```
[26]: # To get a statics information on the numerical column of the data  
df.describe()
```

```
[26]:
```

	stars	review_count	is_open	review_stars	useful	\
count	75517.000000	75517.000000	75517.0	75517.000000	75517.000000	
mean	3.790239	379.438789	1.0	3.794391	1.150800	
std	0.544166	588.797230	0.0	1.379577	3.849485	
min	1.000000	3.000000	1.0	1.000000	0.000000	
25%	3.500000	75.000000	1.0	3.000000	0.000000	
50%	4.000000	171.000000	1.0	4.000000	0.000000	
75%	4.000000	386.000000	1.0	5.000000	1.000000	

max	5.000000	2855.000000	1.0	5.000000	758.000000
-----	----------	-------------	-----	----------	------------

	funny	cool
count	75517.000000	75517.000000
mean	0.402294	0.506826
std	3.224632	2.318385
min	0.000000	0.000000
25%	0.000000	0.000000
50%	0.000000	0.000000
75%	0.000000	0.000000
max	786.000000	321.000000

```
[27]: # Removing the un-wanted columns from the dataset
drop_columns =
↳ ['name', 'address', 'stars', 'review_count', 'is_open', 'attributes', 'categories']
df= df.drop(drop_columns, axis=1)
```

```
[28]: # Creating a new column which gives the number of words in the text column
df['text length'] = df['text'].apply(len)
df.head()
```

```
[28]:
```

	business_id	user_id	review_stars	useful	\
0	RrapAhd8ZxCj-iue7fu9FA	pBUsRjJLTN-TVsoIv5ue8w	3	0	
1	Yy_SB9r7VQTlIIVv_MeOmaA	h1UcaSPIPPqQMiWf12Csqw	4	2	
2	ZDx7kt4aOPTlmYTqXDrTGA	zFGpxwJewI60jC2u9EnZ-g	5	7	
3	RE4qn28MEiDrM1PbdYVgxA	F-nLfSJ4qtdZRwNW993lSw	4	0	
4	_PKXarw3GjlbwbXhjdpuMA	R_TJQ6Hy1BtOYK_hC125bQ	5	1	

	funny	cool	text	text length
0	0	0	Ordered biryani for tonight's supper. I wasn't...	475
1	1	1	We started with a mix platter which had all ki...	323
2	1	2	You know this place is good when most of the p...	1929
3	0	0	Visited this place with my Office colleagues, ...	419
4	1	0	So spicy and flavorful! Shrimp Pakoras really ...	320

```
[29]: # To check if there is any missing data in the new dataFrame
df.isnull()
```

```
[29]:
```

	business_id	user_id	review_stars	useful	funny	cool	text	\
0	False	False	False	False	False	False	False	
1	False	False	False	False	False	False	False	
2	False	False	False	False	False	False	False	
3	False	False	False	False	False	False	False	
4	False	False	False	False	False	False	False	
...	
75512	False	False	False	False	False	False	False	
75513	False	False	False	False	False	False	False	

75514	False	False	False	False	False	False	False
75515	False	False	False	False	False	False	False
75516	False	False	False	False	False	False	False

	text length
0	False
1	False
2	False
3	False
4	False
...	...
75512	False
75513	False
75514	False
75515	False
75516	False


[75517 rows x 8 columns]

12 5.2.2 Data Visualisation

```
[30]: # Importing the libraries to visualise the data
import matplotlib.pyplot as plt
import seaborn as sns
sns.set_style('white')
%matplotlib inline
```

```
[31]: # Using FacetGrid from the seaborn library created a grid of 5 histograms of
↪ text length based on the star ratings
g = sns.FacetGrid(df,col='review_stars')
g.map(plt.hist,'text length')
```


```
[31]: <seaborn.axisgrid.FacetGrid at 0x16b112a54c8>
```



The above graph shows that the more is the review_stars the length of the text lies between 0-1000

```
[47]: # Counting the number of occurrences for each type of star rating
sns.countplot(x='review_stars',data=df,palette='rainbow')
```

```
[47]: <matplotlib.axes._subplots.AxesSubplot at 0x1dff2cca748>
```



The above graph shows that the Indian restaurants have received more positive reviews

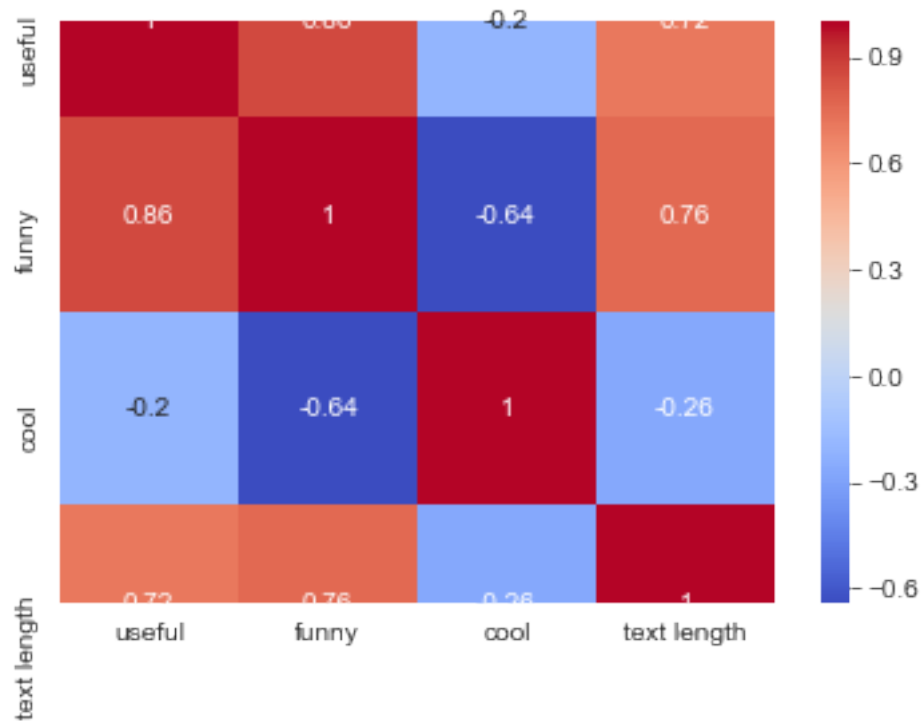
```
[48]: # calculating the mean values of the numerical columns, grouping it by the
      ↪category, stars
stars = df.groupby('review_stars').mean()
stars
```

```
[48]:
```

	useful	funny	cool	text length
review_stars				
1	1.349011	0.624841	0.210544	611.460820
2	1.303633	0.552373	0.319942	686.313710
3	1.216826	0.465370	0.542514	669.506585
4	1.335686	0.447595	0.723777	580.924491
5	1.037770	0.295098	0.501137	435.432785

```
[49]: # Visualising the correlation between the dataframe stars
sns.heatmap(stars.corr(),cmap='coolwarm',annot=True)
```

```
[49]: <matplotlib.axes._subplots.AxesSubplot at 0x1dfea247888>
```



13 5.2.3 Sentiment Detection

```
[37]: # Classifying the dataset and splitting it into the reviews and stars.
# Here, we will classify the dataset into 3 types of rating. Rating 1 = "negative", 3 = "Average", and 5 = "Positive".
data_class = df[(df.review_stars==1) | (df.review_stars==3) | (df.
    review_stars==5)]
```

```
[38]: # Creating the feature and target. x is the 'text' column of data_class and y is the 'stars' column of data_class.
X = data_class['text']
y = data_class['review_stars']
print(X.head())
print(y.head())
```

```
0    Ordered biryani for tonight's supper. I wasn't...
2    You know this place is good when most of the p...
4    So spicy and flavorful! Shrimp Pakoras really ...
5    This is a good take-out place for hakka food i...
6    Th service is really good and the owners are a...
Name: text, dtype: object
0    3
```

```

2    5
4    5
5    3
6    3
Name: review_stars, dtype: int32

```

14 5.2.4 Preparing the data for predictive Modelling

Below I have defined a function that will clean the dataset by removing stop words and punctuations. nltk stands for natural language Toolkit. It is a suite of libraries and programs for symbolic and statistical natural language processing.

```

[99]: import nltk
      from nltk.corpus import stopwords

      # CLEANING THE REVIEWS - REMOVAL OF STOPWORDS AND PUNCTUATION
      def text_process(text):
          nopunc = [char for char in text if char not in string.punctuation]
          nopunc = ''.join(nopunc)
          return [word for word in nopunc.split() if word.lower() not in stopwords.
                  ↪words('english')]

[101]: # Using the fit_transform method on the CountVectorizer object and passing the
      ↪ 'text' column. Saved the result by overwriting x.
      vocab = CountVectorizer(analyzer=text_process).fit(X)
      x = vocab.transform(X)

```

15 5.3 Train Test Split

In order to apply the models, we have to split the data into train and test data. Here, I have divided the data into 70-30. I have taken 70% of the data as a train and remaining 30% of the data as test.

```

[103]: from sklearn.model_selection import train_test_split

[104]: x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.
      ↪ 3, random_state=101)

```

16 5.4 Applying the Machine Learning Models

In this case, we are applying classification learning algorithms. It is a supervised learning process. It is a process where the computer program learns from from the input and then uses this to classify new observations.

The basic idea for opting a classification model is when the target variable is categorical. In this case, we are trying to predict the rating of a review. We have got 5 ratings(0,1,2,3,4,5) out of which we need to predict one. Therefore, classification models can help us in this.

In my analysis, I have used Naive Bayes Classifier, Random Forest Classifier, Decision Trees, K-Nearest Neighbor, RNN Model from Neural Networks.

17 5.4.1 Training a Model using Naive Bayes classifier

Firstly, I want to establish a model that will act as a baseline. In general terms a linear model is appropriate and has the advantage of being fast to train.

Here, I have used Multinomial Naive Bayes over Gaussian because with a sparse data, Gaussian Naive Bayes assumptions aren't met. A simple gaussian fit over the data will not give us a good fit or prediction.

Naive Bayes classification technique is based on Bayes' Theorem with the assumption of independence among predictors. This classifier is easy to build. It is known to be a good fit for very large datasets due to its high scalability.

```
[105]: # Importing the required libraries.
from sklearn.metrics import confusion_matrix, accuracy_score, roc_auc_score, \
    ↪roc_curve
from sklearn.metrics import classification_report
```

```
[106]: # Multinomial Naive Bayes
from sklearn.naive_bayes import MultinomialNB
mnb = MultinomialNB()
mnb.fit(x_train,y_train)
predmnb = mnb.predict(x_test)
print("Confusion Matrix for Multinomial Naive Bayes:")
print(confusion_matrix(y_test,predmnb))
print("Score:",round(accuracy_score(y_test,predmnb)*100,2))
print("Classification Report:",classification_report(y_test,predmnb))
```

Confusion Matrix for Multinomial Naive Bayes:

```
[[2215  377  194]
 [ 339 1445  826]
 [ 101  263 9308]]
```

Score: 86.06

Classification Report:		precision	recall	f1-score	support
1	0.83	0.80	0.81	2786	
3	0.69	0.55	0.62	2610	
5	0.90	0.96	0.93	9672	
accuracy			0.86	15068	
macro avg	0.81	0.77	0.79	15068	

weighted avg	0.85	0.86	0.85	15068
--------------	------	------	------	-------

The performance score of Naive Bayes classifier is 86.06. Since it is high score, I will treat this model as my baseline.

18 5.4.2 Random Forest Classifier

There is no correlation between our feature(text) and target(review_stars) and this is the reason for choosing Random Forest Classifier. The vital thing for a Random Forest Classifier model to make an accurate class prediction is the trees of the forest and more importantly their predictions need to be uncorrelated (or at least have low correlations with each other).

Random forests are an ensemble learning method for classification. It operates by constructing a multitude of decision trees at training time and outputting the class that is the mode of the classes (classification) or mean prediction (regression) of the individual trees.

```
[107]: # Random Forest
from sklearn.ensemble import RandomForestClassifier
rmfr = RandomForestClassifier()
rmfr.fit(x_train,y_train)
predrmfr = rmfr.predict(x_test)
print("Confusion Matrix for Random Forest Classifier:")
print(confusion_matrix(y_test,predrmfr))
print("Score:",round(accuracy_score(y_test,predrmfr)*100,2))
print("Classification Report:",classification_report(y_test,predrmfr))
```

```
C:\Users\sushb\AppData\Local\Continuum\anaconda3\lib\site-
packages\sklearn\ensemble\forest.py:245: FutureWarning: The default value of
n_estimators will change from 10 in version 0.20 to 100 in 0.22.
```

```
"10 in version 0.20 to 100 in 0.22.", FutureWarning)
```

Confusion Matrix for Random Forest Classifier:

```
[[1774  269  743]
 [ 419  690 1501]
 [ 139  259 9274]]
```

Score: 77.9

Classification Report:		precision	recall	f1-score	support
1	0.76	0.64	0.69	2786	
3	0.57	0.26	0.36	2610	
5	0.81	0.96	0.88	9672	
accuracy		0.78		15068	
macro avg	0.71	0.62	0.64	15068	
weighted avg	0.76	0.78	0.75	15068	

The performance score of Random Forest Classifier is 77.9.

19 5.4.3 Decision Tree

Decision Trees are of two types a) classification b) regression. Since the decision variable(target) is categorical/discrete we will be using decision tree classifier. It builds the model in the form of tree structure. The classifier breaks down a data set into smaller subsets and at the same time an associated decision tree is incrementally developed. Decision trees can handle both categorical and numerical data.

```
[108]: # Decision Tree
from sklearn.tree import DecisionTreeClassifier
dt = DecisionTreeClassifier()
dt.fit(x_train,y_train)
preddt = dt.predict(x_test)
print("Confusion Matrix for Decision Tree:")
print(confusion_matrix(y_test,pred dt))
print("Score:",round(accuracy_score(y_test,pred dt)*100,2))
print("Classification Report:",classification_report(y_test,pred dt))
```

Confusion Matrix for Decision Tree:

```
[[1814  510  462]
 [ 481 1079 1050]
 [ 366  804 8502]]
```

Score: 75.62

Classification Report:		precision	recall	f1-score	support
1	0.68	0.65	0.67	2786	
3	0.45	0.41	0.43	2610	
5	0.85	0.88	0.86	9672	
accuracy			0.76	15068	
macro avg	0.66	0.65	0.65	15068	
weighted avg	0.75	0.76	0.75	15068	

The performance score of Decision Tree model is 75.62

20 5.4.4 K Nearest Neighbour Algorithm

KNN is known as a non-parametric and lazy learning algorithm. It is a supervised classification technique that uses proximity as a proxy for 'sameness'. This algorithm takes a bunch of labelled points and uses them to learn how to label other points. To label a new point, it looks at the labelled points closest to that new point (those are its nearest neighbors). Closeness is typically expressed in terms of a dissimilarity function. Once it checks with 'k' number of nearest neighbors, it assigns a label based on whichever label the most of the neighbors have.

```
[109]: from sklearn.neighbors import KNeighborsClassifier
knn = KNeighborsClassifier(n_neighbors=10)
```



```

knn.fit(x_train,y_train)
predknn = knn.predict(x_test)
print("Confusion Matrix for K Neighbors Classifier:")
print(confusion_matrix(y_test,predknn))
print("Score: ",round(accuracy_score(y_test,predknn)*100,2))
print("Classification Report:")
print(classification_report(y_test,predknn))

```

Confusion Matrix for K Neighbors Classifier:

```

[[ 652   60 2074]
 [ 113  137 2360]
 [   56   48 9568]]

```

Score: 68.74

Classification Report:

	precision	recall	f1-score	support
1	0.79	0.23	0.36	2786
3	0.56	0.05	0.10	2610
5	0.68	0.99	0.81	9672
accuracy			0.69	15068
macro avg	0.68	0.43	0.42	15068
weighted avg	0.68	0.69	0.60	15068

The performance score of K Neighbors Classifier is 68.74

21 5.4.5 RNN Model

The RNN is an expressive model that is known to learn highly complex relationships from an arbitrarily long sequence of data. It maintains a vector of activation units for each element in the data sequence, this makes RNN very deep. The depth of RNN leads to two well-known issues, the exploding and the vanishing gradient problems.

There are many ways to implement a neural network in python. Here, I will be using tensorflow/keras.

```

[110]: # Importing the libraries
from keras.preprocessing.text import Tokenizer
from keras.preprocessing.sequence import pad_sequences

```

```

[111]: tk = Tokenizer(lower = True)
tk.fit_on_texts(X)
X_seq = tk.texts_to_sequences(X)
X_pad = pad_sequences(X_seq, maxlen=100, padding='post')

```

```
[112]: from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X_pad, y, test_size = 0.25,
↳random_state = 1)
```

```
[113]: batch_size = 64
X_train1 = X_train[batch_size:]
y_train1 = y_train[batch_size:]
X_valid = X_train[:batch_size]
y_valid = y_train[:batch_size]
```

```
[114]: from keras.models import Sequential
from keras.layers import Embedding, LSTM, Dense, Dropout
vocabulary_size = len(tk.word_counts.keys())+1
max_words = 100
embedding_size = 32
model = Sequential()
model.add(Embedding(vocabulary_size, embedding_size, input_length=max_words))
model.add(LSTM(200))
model.add(Dense(1, activation='sigmoid'))
model.compile(loss='binary_crossentropy', optimizer='adam',
↳metrics=['accuracy'])
```

```
[115]: model.
↳fit(X_train1,y_train1,validation_data=(X_valid,y_valid),batch_size=batch_size,epochs=10)
```

C:\Users\sushb\AppData\Local\Continuum\anaconda3\lib\site-packages\tensorflow_core\python\framework\indexed_slices.py:433: UserWarning: Converting sparse IndexedSlices to a dense Tensor of unknown shape. This may consume a large amount of memory.

"Converting sparse IndexedSlices to a dense Tensor of unknown shape. "

Train on 37604 samples, validate on 64 samples

Epoch 1/10

37604/37604 [=====] - 241s 6ms/step - loss: -211.7108 - accuracy: 0.1830 - val_loss: -331.5008 - val_accuracy: 0.2656

Epoch 2/10

37604/37604 [=====] - 253s 7ms/step - loss: -548.4435 - accuracy: 0.1831 - val_loss: -620.0129 - val_accuracy: 0.2656

Epoch 3/10

37604/37604 [=====] - 253s 7ms/step - loss: -880.8150 - accuracy: 0.1831 - val_loss: -908.4469 - val_accuracy: 0.2656

Epoch 4/10

37604/37604 [=====] - 253s 7ms/step - loss: -1213.9509 - accuracy: 0.1831 - val_loss: -1196.3496 - val_accuracy: 0.2656

Epoch 5/10

37604/37604 [=====] - 253s 7ms/step - loss: -1545.9127 - accuracy: 0.1831 - val_loss: -1483.9255 - val_accuracy: 0.2656

Epoch 6/10

```

37604/37604 [=====] - 213s 6ms/step - loss: -1877.2001
- accuracy: 0.1831 - val_loss: -1770.8112 - val_accuracy: 0.2656
Epoch 7/10
37604/37604 [=====] - 173s 5ms/step - loss: -2208.1054
- accuracy: 0.1831 - val_loss: -2057.5815 - val_accuracy: 0.2656
Epoch 8/10
37604/37604 [=====] - 171s 5ms/step - loss: -2538.9139
- accuracy: 0.1831 - val_loss: -2344.1816 - val_accuracy: 0.2656
Epoch 9/10
37604/37604 [=====] - 171s 5ms/step - loss: -2869.9704
- accuracy: 0.1831 - val_loss: -2632.5930 - val_accuracy: 0.2656
Epoch 10/10
37604/37604 [=====] - 171s 5ms/step - loss: -3204.0844
- accuracy: 0.1831 - val_loss: -2921.8779 - val_accuracy: 0.2656

```

[115]: <keras.callbacks.callbacks.History at 0x16b2363ba08>

The performance score of RNN Model is 26.56

22 6. Score of the above classifier models-- Results

Multinomial Naive Bayes -- 86.06, Random Forest Classifier -- 77.9, Decision Tree -- 75.62, K Nearest Neighbour Classifier -- 68.74, RNN -- 26.56.

Multinomial Naive Bayes has the best score. We will use this model to predict a random positive, negative and average review.

23 7. Validation

24 7.1 Predict positive review

```

[134]: # POSITIVE REVIEW
pre = df['text'][20]
print(pre)
print("Actual Rating: ",df['review_stars'][20])
pre_t = vocab.transform([pre])
print("Predicted Rating:")
mnb.predict(pre_t)[0]

```

Unexpectedly wonderful Indian cuisine, I had the Tandoori Mixed Grill. Huge portion with a mix of plump juicy shrimp, amazingly tender chicken and delicious fish, tons of vegetables and great flavors! They served rice with Tikka Masala and another garlic sauce, just tremendous. I ordered hot and it was not too spicy, next time I might order the next level hotter, but it was excellent in any case. Very attentive service, too!

Actual Rating: 5
Predicted Rating:

[134]: 5

25 7.2 Predict Average Review

```
[235]: # AVERAGE REVIEW
ar = df['text'][6]
print(ar)
print("Actual Rating: ",df['review_stars'][6])
ar_t = vocab.transform([ar])
print("Predicted Rating:")
mnf.predict(ar_t)[0]
```

Th service is really good and the owners are awesome. Very friendly and seems like a family run establishment.

I thought the food was mediocre though not as good as federicks in Markham but maybe I ordered the wrong thing.

Will give it another try. Not bad if you're in the area for lunch looking for cheap eats. The \$10 lunch special is well worth the money.

Actual Rating: 3
Predicted Rating:

[235]: 3

26 7.3 Predict Negative Review

```
[229]: # NEGATIVE REVIEW
nr = df['text'][58]
print(nr)
print("Actual Rating: ",df['review_stars'][10])
nr_t = vocab.transform([nr])
print("Predicted Rating:")
mnf.predict(nr_t)[0]
```

The more I remember this, the worse it gets. This is the second worst indian food I've ever had in my life (first being Maezo). We had takeout: The lamb roganjosh was very bland, and the meat did not even taste like lamb. It looked like chicken. I think they messed up our order.

The chicken masala was terrible compared to other chicken masalas I've had at

other restaurants. It was overly oily, and the peppers in it tasted rancid.

The naan was wet and soggy. They didn't pack it properly to prevent the steam collecting inside the aluminum. Other restaurants are able to pack it to-go without making the naan soggy. The tandoor paneer was actually alright, it was the only acceptable part of our meal.

I later told my Indian friend about this, and he said "oh Aroma? why on earth would you go there?" Well, now I know...

Actual Rating: 4

Predicted Rating:

[229]: 1

27 8. Summary:

From the datasets, we have found that: 1 -- Mon Ami Gabi is the merchant who has got the maximum number of positive reviews. 2 -- The category of top most reviewed business is restaurants. 3 -- Casinos have got the most negative reviews.

From the machine learning models for sentiment analysis, it is clear that Multinomial Naive Bayes performs the best. The validation proves that the model works fine for positive and average reviews. But it seems to be not working for the negative reviews.

28 9. Future Scope:

I believe the reason for which the model fails to predict the negative review is that there are more positive reviews as compared to the negative ones in the dataset(the collected data,df). That means the dataset is not normally distributed. I can suggest two ways to improve it:

1 -- We can normalize the data. So that the positive and negative ratings are equally distributed over the dataset.

2 -- While collecting the data, we can also check with other business categories (Shopping, food, home services,etc). It might be possible that only the reviews for "Indian Restaurants" contain mostly positive ones.

[]: